Revenue Workspace — Product Requirements v3.1

Status: current. **Date:** 13 October 2026; **v3.1 update:** 15 October 2026 (see Changelog). **Supersedes:** PRD v2 (12 October 2026) and PRD v1 (6 October 2026). **Inputs:** the product owner's answers of 13 October 2026, the Lara meeting transcript, and eleven research reports in `docs/research/` (report 11, added in v3.1, is an anonymised teardown of a real outbound engagement). **Current authorization:** local implementation, focused commits and project-only Azure deployment and inference (owner, 15 October); no repository push or live outreach.

v3 records the product owner's decisions on the v2 proposals and adds the business model, the "one place" component plan, the out-of-the-box plays and the Asset Studio. This document is about the product; competition rules live in research 06. Every feature from v1 and v2 stays in the catalogue (§7 and Appendix A) so the team can pull work forward if it finishes early.

0. Summary

Revenue Workspace is **one place where non-technical go-to-market (GTM) people run outbound end to end**: find the right accounts, research them, write the first messages, make the assets that open conversations, run the sequence and keep every result in a CRM. Today a skilled "GTM engineer" builds this alone from 6–8 paid tools wired to Claude, Clay or n8n with personal connectors. Most founders, sales people and small agencies cannot do that: they do not clone GitHub repos or wire APIs. Many of the parts already exist, but they are scattered. We **gather them, tune them and deliver them in one product**, with three things on top:

1. **Plays out of the box.** Ready-made, versioned workflows ("first campaign in one chat", "funding-signal outreach", "mini-audit outreach", "weekly client report"). A user starts from a play, not from a blank canvas.
2. **An AI assistant that runs the plays.** The user describes the goal in chat; the assistant sets up the play, shows the cost before spending, drafts, and asks for approval before anything is sent.
3. **An Asset Studio built on OpenDesign.** Branded one-pagers, mini-audits, case studies, lead magnets, carousels, proposals and client reports, in English and Arabic, generated from CRM data and attached to the prospect record.

Underneath sit the parts that personal automations lack: team connectors and vault, spend caps, approvals, suppression, audit, and Egypt/Saudi/UAE compliance gates.

Language: English is the default for the interface and content. Arabic content (Egyptian first, Gulf as beta, business Arabic) is a first-class option and the main advantage when we pitch in MENA later. **Business model:** freemium, like Supabase: a useful free tier with hard limits, then Solo, Team, Agency and Enterprise plans with pooled credits and spend caps. **Cloud:** Azure is a placeholder because it is the only cloud where we have credits. We use it to test the demo with real services. Huawei Cloud comes only if we pass the December phase; then we migrate. Every provider-specific service sits behind an adapter so the move is configuration, not code. **Source:** the code will be open source.

1. Decision log (settled 13 October 2026)

All decisions below are settled by the product owner. Do not reopen a settled decision without new facts.

#	Topic	Decision	Source / reason
D1	Language	English is the main language: default UI, default generated content, competition material. Arabic content (Egyptian, Gulf beta, business Arabic) is supported in drafts and assets. Arabic UI later through Twenty's native RTL	Owner. Judges are international; Arabic is the MENA startup pitch. GoHighLevel's Arabic request has 308 votes open since 2023 (research 10)
D2	Users and business model	Freemium. Free tier for solopreneurs with hard limits; paid tiers for founder-led teams (1–10) and small agencies. Tier table in §4	Owner ("learn the Supabase model"). Research 07
D3	Scope	Keep the full feature catalogue (§7, Appendix A). Build the outbound loop first; pull later features forward if time allows	Owner
D4	Cloud	**Azure is a placeholder** for the December demo, so the demo runs on real services; we have credits only there. If we pass the December phase, we receive Huawei Cloud resources and migrate (§8.3). The December submission is a PDF only; the demo is a bonus and we do not submit the repository	Owner, 13 Oct
D5	Thesis	"Gather, tune and deliver in one place" for non-technical GTM users: plays + assistant + Asset Studio + governed connectors	Owner. Research 08 (8 tools per team, 42% overwhelmed), research 10 (learning curve and setup time are the top complaints)
D6	Asset Studio	OpenDesign is a major pillar, not a side feature. Ship a template library (§5.4)	Owner. Research 09 (OpenDesign 0.24.1, Apache-2.0, 116 templates, brand-from-URL)
D7	LinkedIn	In the product: manual LinkedIn tasks with copy-ready text, own-profile posting (`w_member_social`), licensed data providers and customer CSV exports. No crawling in the product or the demo. A lower-risk method for deeper LinkedIn data is decided when we reach it (§7.7, §9 item 1)	Owner, 13 Oct. Research 09
D8	WhatsApp	Opt-in only, after a reply or a form; official Cloud API; never cold	Settled in the meeting
D9	Autonomy	AI researches and drafts; a human approves; a sender platform sends at capped volume. No autonomous AI SDR	Settled. Research 01 (11x case)
D10	Compliance gates	Lawful basis, country channel rules, opt-out, suppression and residency flags are in the product	Owner confirmed
D11	Credits and keys	Hybrid: each paid plan includes pooled AI and enrichment credits; bring-your-own-key (BYOK) for data providers removes the data-credit charge. Tier prices in §4.2 are used now and hardened with the business model later	Owner, 13 Oct. Research 07
D12	Competition data	Synthetic or consented data only in every competition artifact	Settled
D13	Assistant	Spike Twenty's native AI chat, skills and MCP before building a custom shell	Settled from v2

#	Topic	Decision	Source / reason
D14	Evidence	Build the business case from public market data (§3). Do not quote unverified figures (for example "6,000 / 130 GTM engineers")	Owner: no follow-ups with Lara; use internet research
D15	Huawei technology	Huawei services are the post-December target: ModelArts Studio (MaaS) for the model, ECS/CCE, RDS, DCS and OBS (§8.3). Not used before December	Owner, 13 Oct
D16	Team eligibility	Confirmed by the owner. Not an open question	Owner
D17	Open source	The project is open source: the code is public on GitHub; revenue comes from the hosted service, credits and agency features (open core, like Supabase and Twenty). Our licence (Apache-2.0 or AGPL) is chosen before the repository goes public. Keep secrets and customer data out of the repository from day one	Owner, 13 Oct
D18 *(v3.1)*	Design system	The product UI uses a Raycast-derived design system (copied from Refero Styles into `design-systems/raycast/`, adapted in `design-systems/revenue/DESIGN.md`): dark default plus light theme, one coral accent reserved for AI, Arabic UI strings and RTL	Owner, 15 Oct
D19 *(v3.1)*	AI harness	Our own agent harness on Vercel AI SDK Core with Azure OpenAI: we own the loop, step cap, prompts (versioned files) and approval interrupts; spending tools never execute without a `SpendApproval`. The provider is swappable to Huawei ModelArts MaaS (OpenAI-compatible)	Owner request for full control, 15 Oct
D20 *(v3.1)*	Azure models and voice	Azure credits confirmed by Amr; provisioning and real inference approved for one project-only Azure OpenAI resource (chat model + Whisper). Voice dictation uses Whisper in English and Arabic	Owner, 15 Oct
D21 *(v3.1)*	Real data	Real customer exports are read for structure only, are git-ignored, and never appear in the repository; demo data is generated synthetically with a leak test	Owner, 15 Oct; extends D12

2. Delivery context

The product is built for the market first; the Huawei ICT Competition (Innovation) is the first milestone. What matters for the product:

- **December:** we send a PDF of the idea (problem, users, solution, business model, market, moat, roadmap). An architecture plan and a demo website are bonuses; we want both, and the demo runs on real services on Azure. We do not submit the repository.
- **If we pass:** we receive Huawei Cloud resources, migrate (§8.3) and build phase **N** toward the regional stage.
- **What judges reward:** assume the 2023–24 rubric unless a newer one is published: Creativity 30, Technical complexity 25, Social value 15, Functional completeness 15, Presentation 15. For the product this means: lead with what is new (plays for non-technical users, Arabic content, MENA compliance), show real technical depth (scoring, cost-controlled waterfall, governed agent tools, Arabic evaluation), state the social value (Arabic-

speaking small businesses get a sales capability that today needs an English-first expert stack), and make every shown flow complete.
- **Evidence:** every number in the PDF comes from the "safe numbers" list in research 08 §7, with the publisher named. Data in every artifact is synthetic or consented.

3. Problem and evidence

Use these numbers in competition material only as written here, with the caution. Full list: research 08 §7.

Claim	Number	Source	Caution
Sales reps spend most time not selling	40% of the week selling, 60% on other tasks	Salesforce State of Sales, 7th ed. (n = 4,050, 22 countries, Aug–Sep 2025)	Vendor survey; no Middle East sample
Tool sprawl	Teams use 8 tools on average; 42% of reps feel overwhelmed by tool count	Same	Same
Demand for one place	84% of teams without an all-in-one platform plan to consolidate	Same	Incumbents see the same demand
Skills gap	Only 43% of sales and marketing teams in large US firms have formal AI training	Zapier / Centiment, Mar 2026	Vendor; large US firms only. SMB gap must be confirmed by interviews
MENA software demand	MENA software spending USD 20.4B in 2026 (+13.9%)	Gartner forecast via Arab News	Quote as "Gartner forecast"
SMEs in the target markets	KSA: about 1.7M active SME commercial records, 8.4M employees (Monsha'at 2025). UAE: more than 1.2M companies, about 95% SMEs (Ministry of Economy, Sep 2025). Egypt: 96.3M internet users	Government sources	Records are not firms
AI familiarity	45.2% of Saudi internet users use AI tools	CST Saudi Internet Report 2025	Consumer use
Arabic gap	Arabic is 0.6% of websites	W3Techs, 12 Oct 2026	Site language only; no demand data for Arabic business tools
AI SDR market	USD 4.1B (2025) to 15.0B (2030)	MarketsandMarkets, Jul 2025	Analyst forecast; publishers differ widely

User pain (research 10, public reviews and forums; Reddit and G2 were not reachable): the top jobs are (1) get a first campaign live without learning a tool, (2) know the cost before clicking and never lose paid credits, (3) keep automations working and be told in plain words when they break, (4) set up domains and inboxes safely. Arabic and RTL support is a medium-strength need. Design assets are a plausible need with **no independent evidence of reply lift**, so every asset play ships with a holdout (§9 item 3).

Outbound reality (research 01): typical cold email reply rate is about 1%; top decile 2.6–3.9%. Signal and personalization "lift" figures are vendor claims.

4. Users and business model

4.1 Segments

Segment	Plan	Why
Solopreneurs and first-time founders	Free, then Solo	Strong pain, lowest budget, high churn; free tier is the acquisition funnel
Founder-led B2B teams, 1–10 people, Egypt/KSA/UAE	Solo or Team	Founder still sells; no playbook; cannot afford a $260–400/month tool stack
Small B2B agencies and lead-gen freelancers, 2–20 people	Agency	Multi-client work, reporting and handover are documented pains (research 10 §5)
Accelerators and VCs	Channel (startup discount code)	Access to founder cohorts
SMEs with 10–50 people	Team / Enterprise later	Bigger budgets, longer cycles

4.2 Tiers (in use now; hardened later with the business model and a real cost log)

Modelled on Supabase (hard-limited free tier, base fee plus usage, spend cap on by default) and on GTM tools that combine seats with a pooled credit meter. Details and cost reasoning: research 07 §6.

	Free	Solo	Team	Agency	Enterprise
Price (USD/month)	$0	$24 monthly / $19 annual	$39 per seat monthly / $32 annual, 3–10 seats	$199 monthly / $159 annual base (test $149–249)	From about $1,000
Seats	1	1	3–10 + free view-only	5 + free client viewers	10+, SSO
Workspaces	1	1	1	5 client workspaces; +$25 each	Custom
Records	1,000	10,000	50,000	25,000 per client	Custom
AI credits / month	100	400	600 per seat, pooled	5,000 pooled	Custom
Enrichment credits / month	25 (BYOK unlimited)	100 (BYOK unlimited)	150 per seat, pooled	1,500 pooled	Custom
Sending	1 sequence, 100 sends/month from own mailbox	5 sequences, 2,000 sends	20 sequences per seat, 5,000 sends per seat	10,000 sends per client	Custom
Plays	Starter plays	All plays	All plays + team plays	All + client templates ("snapshots")	Custom plays
Asset Studio exports / month	5, with "Made with" mark	50	200 per seat, shared brand kit	200 per client, client brand kits	Custom

	Free	Solo	Team	Agency	Enterprise
Arabic and English	Yes	Yes	Yes	Yes	Yes
Controls	Hard stop at quota; archive after 60 days idle	Spend cap on by default; prepaid packs	Roles, per-member caps, approvals, audit	Per-client caps, client rebilling, white-label reports	SSO, SLA, residency, invoicing

Rules: do not gate Arabic; gate volume. Show the credit cost of each action before it runs. Credits reset monthly; prepaid packs cost more and do not expire. No cliff between tiers (each step under 2.5×). Startup code: 50% off Solo or Team for 6 months for partner accelerators.

Payments: Stripe lists the UAE but not Egypt or Saudi Arabia (research 07, verified). For the competition: no real billing, plans and a mock checkout only. First paying users: a merchant of record (Paddle first; confirm in writing that an Egypt- or KSA-based seller is accepted) or a UAE entity with Stripe. Later: Paymob (EGP) and Moyasar or HyperPay (SAR, mada recurring). VAT (KSA 15%, UAE 5%, Egypt 14%) needs a tax adviser before launch.

5. Product: one place

5.1 Pillars

Pillar	What the user sees	What we gather and tune	What we build
Plays	A library of ready workflows; "start from a play"	Patterns from Clay templates, lemlist playbooks, GoHighLevel snapshots	Play engine, versioning, MENA templates, holdouts
Assistant	Chat that sets up and runs plays, explains scores and sources, asks for approval	Twenty AI chat, skills and MCP (spike first)	Task-level product tools, approval and cost gates
Asset Studio	Branded assets from CRM data, edited in chat, saved on the record	OpenDesign engine, templates, brand-from-URL	Per-tenant wrapper, our outreach templates, Arabic layouts, PDF/PNG render
Data and signals	Accounts, contacts and signals with honest coverage	Licensed providers (BYOK), PredictLeads/Crustdata, Crawl4AI for public web	Waterfall, caching, cost caps, coverage report
Sending	Sequences across email, LinkedIn tasks and call tasks	Instantly or Smartlead for email; customer's own mailboxes	Sequence state, reconciliation, deliverability pre-flight
CRM	One record of everything	Twenty CRM (as an app on its API/SDK; no fork)	Product objects, play records, audit
Governance	Team vault, budgets, approvals, suppression, compliance	—	All of it (§7.9–7.10)

5.2 Out-of-the-box plays

Every play: human approval before sending, cost shown before spending, English and Arabic from the same play, a plain-language run log, results measured as replies per delivered message, holdout option on every asset play. Source: research 10 §7(b).

#	Play	Trigger	Output	Phase
P1	First campaign in one chat	New workspace	ICP card, offer one-pager, 3-step sequence	D
P2	Inbox and domain pre-flight	Mailbox or domain connected	Health report, caps, warm-up calendar, launch block until checks pass	D
P3	Funding or hiring signal outreach	New funding round or job post at a target account	Signal brief, "idea for you" one-pager, drafts	D
P4	Mini-audit outreach	List with websites (strong for Egyptian and Saudi SMEs)	One-page audit (en/ar), personal page link, drafts	D
P5	Lead magnet in a day	User picks an offer	Checklist or short report PDF, gated page, 3 posts	D+
P6	Case study builder	Deal won or happy-client message	Case study one-pager, carousel, post	D+
P7	Weekly LinkedIn carousel	Weekly schedule	Carousel PDF and caption (posted by the user or via `w_member_social`)	N
P8	Opt-in WhatsApp follow-up	Contact from a meeting, event or form	Consent log, approved template message	N
P9	Reply handler	Inbound reply	Classification, draft reply, record update	D+
P10	Meeting-booked proposal	Meeting scheduled	Research brief, one-page proposal, confirmation email	N
P11	Revive cold leads	60–90 days without reply	New angle, short asset, one touch, then stop	N
P12	Agency client onboarding kit	New client	Intake, brand kit from URL, client workspace from template, guided mailbox connection	N
P13	Weekly client report	Weekly	Sends, replies, meetings (no opens); plain-language note; branded PDF/link	N
P14	Missing-data fallback (GCC)	Lead without email or phone	Provider order, coverage score, manual LinkedIn/WhatsApp task	D+
P15 *(v3.1)*	Discovery pilot	New segment or pre-PMF offer	Interview guide (en/ar), questionnaire, capture form with pain scores and maturity 0–4, launch gate, scorecard	D+
P16 *(v3.1)*	Event list follow-up	Attendee or exhibitor list after an event	Segments by register, company-level event reference rule, multi-channel cadence, HOLD for missing research	D+
P17 *(v3.1)*	Partnership campaign	Partner-bucket accounts (channel, accelerators, consultancies)	Partner angle, Day 0/2/7/16 cadence, connect → accepted → reply → meeting → agreed scorecard	N

Plays from the meeting's process (§6) are the engine underneath: segment, signals, scoring rubric, personas and offers, enrichment policy, research brief, message framework, sequence and guardrails are parts of every play.

5.3 Assistant behaviour

The assistant follows the v1 action contract (v1 §4: context, action levels, typed tools, durable jobs). It calls only product tools (`find_signals`, `score_accounts`, `enrich_contacts`, `draft_messages`, `make_asset`, `queue_sequence`), never raw provider tools. It explains every score and fact with its source, shows the cost estimate before any spend, and stops for approval at the defined action levels. Retrieved web text is untrusted and cannot trigger tools.

5.4 Asset Studio (OpenDesign)

OpenDesign 0.24.1 (Apache-2.0) is a single-user local daemon with an HTTP API, a stdio MCP server, 116 templates, 151 design systems and brand-from-URL. Exports: HTML, PDF, PPTX, ZIP, Markdown, MP4 (no PNG). It has no multi-tenant mode (research 09 §B10). We run it as an internal service behind our own auth, one workspace volume per tenant, through its BYOK model path so no coding-agent CLI is spawned. Our templates are added as OpenDesign skills, not as a fork. PNG and A4 PDF are rendered by headless browser.

Templates to ship, in order (research 10 §7(c)): offer one-pager; mini-audit report; case study (page + carousel); checklist lead magnet with gated page; LinkedIn carousel (1080×1350); per-prospect landing page with view tracking; short proposal; WhatsApp card; weekly client report; email header image. Rules: one brand kit per workspace or client, pulled from the website; English and Arabic are two real layouts (RTL tested), not a mirrored copy; every figure on an asset links to a source; the asset is stored on the prospect or account record; the user can change one element without regenerating the whole asset.

6. The process we productize (from the meeting)

Lara's outbound process, with transcript timestamps, as recorded in v2:

#	Stage	What the GTM person does today	Tools named	Pain
1	ICP	Market research; closest competitor's announced clients; 3–4 segments (00:05:14–00:06:26)	ChatGPT/Claude deep research	Most firms have no ICP book or playbook (00:07:52)
2	Offer per persona	Same service, different entry per role (00:00:22–00:01:45)	Docs	Offers not written down per persona
3	Signals	Hiring, funding, headcount growth, job change, keyword posts (00:06:42–00:07:46)	Crunchbase, Crustdata, Clay, PhantomBuster, Apify, Bright Data	Ready tools miss custom signals; scraping is imprecise
4	Scoring	Score out of 100; under 50 disqualified; A/B/C tiers (00:16:34–00:17:21)	Spreadsheets, Clay	Done by hand
5	Waterfall enrichment	Cheapest provider first until verified (00:17:28–00:19:17)	Apollo, Lusha, FullEnrich, Surfe, Clay, Mawsool	Poor GCC coverage; phones expensive
6	Research per prospect	A recent, specific fact to open with (00:19:21–00:20:15)	LLM research	Generic messages get no replies
7	Lead magnet	Free useful asset that proves quality (00:20:58–00:21:49)	Docs, design tools	Trust-heavy markets resist first touches

#	Stage	What the GTM person does today	Tools named	Pain
8	Multichannel sequence	Email → LinkedIn invite → email → LinkedIn → call (00:03:50–00:04:44)	Smartlead/Instantly, HeyReach/Dripify	Domain blocks; warm-up and mailbox setup
9	CRM and learning	Record replies and next steps	HubSpot, Zoho, GoHighLevel	Data spread across tools
—	Glue	Each tool wired to Claude/Claude Code with connectors and scheduled tasks (00:15:24–00:15:44)	Claude, MCP	Every person builds their own

7. Requirements

Phase codes: **D** = December demo website (Azure). **D+** = stretch before December. **N** = only if we pass December (Huawei migration and the regional build). **L** = later (startup roadmap). Requirement IDs from v2 are kept; new v3 IDs are marked *(new)*. Appendix A lists every v1 feature with its v3 phase.

7.1 Onboarding and plays

ID	Requirement	Phase	Acceptance
ONB-01 *(new)*	Website-URL onboarding: AI drafts company context, ICP, offers and brand kit with cited sources	D	Draft separates facts (cited) from hypotheses; user edits before saving; brand kit applies to the first asset
ONB-02 *(new)*	Guided connections in plain language (mailbox OAuth, BYOK provider keys, sender platform) with a health check	D	A non-technical tester connects a mailbox and one provider in under 10 minutes without documentation
PLAY-01 *(new)*	Play library with the plays in §5.2; start from a play	D (P1–P4), D+ (P5, P6, P9, P14), N (rest)	A new workspace reaches an approved draft from a play in ≤ 30 minutes
PB-01	ICP builder: segments with firmographic filters stored as Twenty records	D	Segment saves, versions and shows which plays use it
PB-03	Offer variants per persona with approved claims and proof	D	A draft cannot use a claim outside the approved list (server-side check)
PB-04	Play builder with versioning; running work keeps its version	D	A new version does not silently change approved drafts
PB-06	Playbook export (PDF) — the sales playbook most firms lack	D+	Export contains ICP, personas, offers, rubric and sequence
PLAY-02 *(new)*	Plain-language run log and failure alerts ("this step failed because…, do this")	D	Every failed step shows cause and next action; no silent failure
PLAY-03 *(new)*	Scheduled plays replace personal scheduled tasks	D	A weekly run appears in the review queue without user action and survives restart

7.2 Signals and sourcing

ID	Requirement	Phase	Acceptance
SIG-01	CSV/XLSX import (including Sales Navigator and LinkedIn connection exports by the customer) with row-level outcomes	D	Every row imported, matched, excluded, rejected or held; re-import does not duplicate
SIG-02	Company signals through one provider (PredictLeads or Crustdata) behind a connector	D	Each signal stores provider, date, URL and confidence
SIG-03	Custom keyword signals from public sources (news, careers pages, job boards) via Crawl4AI	D+	Match stores source URL and excerpt
SIG-04	Funding signals (Crunchbase or a regional tracker)	N	Round, date, amount and source; missing data stays unknown
SIG-05	Website mini-audit checks for play P4 (speed, mobile, contact paths, Arabic version, social proof)	D	Each check stores the evidence it used

7.3 Scoring and enrichment

ID	Requirement	Phase	Acceptance
SCR-01	Deterministic rubric scoring (0–100) with per-rule explanation	D	Same input, same score; missing field shows "unknown"
SCR-02	Tier cut-offs; disqualification before enrichment spend	D	Disqualified accounts never reach enrichment
SCR-03	AI-suggested rule changes from won/lost outcomes	L	Shows the data; activation needs approval
ENR-01	Waterfall: cache/CRM → provider 1 → 2 → 3; stop on verified result, provider limit or cost cap	D	Order, stop reason and cost per field logged
ENR-02	Pre-flight cost estimate; caps per run, per member, per workspace per month	D	A run above the cap does not start; estimate vs actual shown
ENR-03	Caching (emails 90 days, phones 30, misses 14)	D	Repeat request inside TTL costs nothing
ENR-04	Phones only for tier A by default	D	Enforced server-side
ENR-05	MENA coverage report per provider and country	N	Built from ≥ 200 test records; shows failures
ENR-06	Fixture providers with the same interface for the demo	D	Fixture results labelled simulated everywhere
ENR-07 *(new)*	Email verification with a paid verifier first and Reacher (MIT, self-hosted) as fallback	D+	Verification result and source stored

7.4 Research, messages and Arabic

ID	Requirement	Phase	Acceptance
RES-01	Per-prospect research brief with dated citations; facts, hypotheses, unknowns separated	D	An uncited "fact" blocks approval
RES-02	Public company sources and user-provided data only	D	Banned sources are never called
RES-03	Retrieved text is untrusted	D	Prompt-injection cases do not change actions
MSG-01	Drafts per channel (email, LinkedIn note, call script) from the play's framework	D	Length and structure rules per channel respected
MSG-02 *(new)*	Voice and tone profile from the customer's own website and posts; few-shot examples with recorded origin	D+	No scraped third-party posts in the example store
ARB-01	Language per contact: English (default), Egyptian, Gulf (beta), business Arabic	D	Field on the Twenty person record; change re-drafts after confirmation
ARB-02	Dialect-conditioned generation with few-shot examples	D	Each example has an origin
ARB-03	Automatic checks: dialect consistency, MSA leakage, English switching, bidi rendering, banned phrases, opt-out line	D+	Failing check shows a reason; reviewer can override with a note
ARB-04	Every Arabic message needs human approval	D	No path sends an unapproved Arabic message
ARB-05	Arabic evaluation harness (rubric, golden set, native reviewers, calibrated LLM pre-filter)	N	Results per dialect, with and without few-shot
ARB-06	Gulf output labelled beta until reviewed by native Gulf reviewers	N	Unreviewed Gulf drafts need an explicit override
ARB-07	Arabic UI via Twenty RTL and translated strings	L	Core pages mirror correctly

7.5 Sequences and sending

ID	Requirement	Phase	Acceptance
SEQ-01	Steps, delays, stop rules; reply or opt-out stops later touches across all plays	D	Fixture reply cancels pending steps
SEQ-02	Email through one sender platform (Instantly or Smartlead); we do not build mailbox infrastructure	D (sandbox /fixture)	Submitted, accepted, sent, failed, replied are separate reconciled states
SEQ-03	Deliverability pre-flight (play P2): SPF/DKIM/DMARC, one-click unsubscribe, bounce < 2%, complaints < 0.1%, ~30 sends per mailbox per day, warm-up age	D	A failing check blocks launch
SEQ-04	LinkedIn steps as manual tasks: copy-ready text, profile link, done/skip	D	No automated LinkedIn action in the product

ID	Requirement	Phase	Acceptance
SEQ-05	Call steps as tasks with script and allowed hours per country	D	Out-of-hours tasks flagged
SEQ-06	Reply inbox and classification (play P9) with human confirmation	D+	Opt-out applies immediately
SEQ-07	WhatsApp only to opted-in contacts via the official Cloud API (play P8)	N	Send blocked without opt-in and approved template
SEQ-08	Holdout groups per play and per asset (asset vs no asset)	D+	Reply-per-delivered with sample size; no lift claim below minimum sample

7.6 Asset Studio

ID	Requirement	Phase	Acceptance
STU-01 *(new)*	OpenDesign service wrapper: own auth, one volume per tenant, BYOK model path, pinned release, no CLI spawn	D	Tenant A cannot read tenant B's projects; engine version recorded per asset
STU-02 *(new)*	Brand kit from website URL (logo, colours, fonts) as a per-workspace `DESIGN.md`	D	First asset uses the kit; user can correct it
STU-03 *(new)*	Templates 1–2 (offer one-pager, mini-audit) in English and Arabic	D	Every figure links to a source; Arabic PDF renders RTL correctly
STU-04 *(new)*	Templates 3–6 (case study, checklist lead magnet + gated page, LinkedIn carousel, per-prospect page with view tracking)	D+ (3–4), N (5–6)	Same as STU-03
STU-05 *(new)*	Templates 7–10 (proposal, WhatsApp card, client report, email header)	N	Same as STU-03
STU-06 *(new)*	Edit in chat: change one element without regenerating the rest; versions and lineage; changed facts reopen review	D+	Revision creates a new version; approved original identifiable
STU-07 *(new)*	Export A4/Letter PDF, PNG (headless render), PPTX where OpenDesign supports it	D (PDF), D+ (PNG), N (PPTX)	Export matches preview
STU-08 *(new)*	Asset stored on the prospect/account record and linked to the play and message	D	Asset opens from the record
DES-02	Account research brief export	D	Matches the approved research record

7.7 LinkedIn access

ID	Requirement	Phase	Acceptance
LI-01 *(new)*	Manual LinkedIn tasks with copy-ready text (SEQ-04)	D	No automated action
LI-02 *(new)*	Customer-uploaded LinkedIn data (connections export, Sales Navigator CSV) through SIG-01	D	Source labelled; consent basis recorded
LI-03 *(new)*	Own-profile posting via official `w_member_social` (plays P6, P7)	N	Founder's OAuth; 150 requests/member/day respected
LI-04 *(new)*	Licensed people data (Crustdata, Apollo, etc.) through BYOK	D+	Provider and date stored per field
LI-05 *(new)*	Local power-user connector (Agent-Reach / `linkedin-mcp-server`): runs only on the user's computer with the user's own session; one profile per click; no search harvesting, no background crawl, no automated messages; warning screen; daily cap; kill switch; pinned version	L — the lower-risk method is chosen when we need deeper LinkedIn data; licensed providers (LI-04) first	Cookie never leaves the user's machine; only user-chosen fields are saved, with source and deletion path
LI-06 *(new)*	Internal team use of Agent-Reach for non-LinkedIn research (web, RSS, YouTube, GitHub) on team laptops	Now (internal)	Not part of the product or the demo

7.8 Assistant

ID	Requirement	Phase	Acceptance
AST-01	Spike: Twenty native AI chat + skills + MCP with our Azure model endpoint vs a custom persistent popup	D (first week)	Written decision covering context, navigation, typed actions, custom LLM endpoint
AST-02	Assistant runs plays step by step, explains scores and sources, asks for approval	D	Full play run from chat; no send without approval
AST-03	Cost shown before every spending action	D	Estimate appears before confirmation
AST-04 *(new)*	"Make me an asset" in chat calls the Asset Studio and returns a preview	D	Preview and record link appear in the chat

7.9 Governance and team

ID	Requirement	Phase	Acceptance
GOV-01	Connector registry: provider, auth, owner, quotas, cost model, risk label, allowed actions	D	Admin sees status, last success/error, remaining quota
GOV-02	Per-workspace vault; BYOK; keys never reach the model, browser or logs	D	Secret scan finds no key in prompts, logs or exports
GOV-03	One product MCP server with task-level tools	D	Every tool call checks role, budget and approval
GOV-04	Rate limiting per provider and credential	D	Burst tests stay inside provider limits
GOV-05	One contact, one owner, one active sequence across the workspace	D	Second enrolment needs a handoff
GOV-06	Roles: owner, operator, reviewer, viewer	D	Restricted roles cannot approve sends or change budgets
GOV-07	Audit log of tool calls, approvals, costs and sends	D	A send traces to play version, draft version, approver and cost
GOV-08	Agency mode: client workspaces, per-client brand kit, budget, playbook, free client viewers, client templates	N	Data and assistant history never cross clients
GOV-09 *(new)*	LLM tracing and cost per tenant with Langfuse (self-hosted, MIT core)	D+	Each run records model, prompt version, tokens, cost

7.10 Compliance (Egypt, KSA, UAE, EU)

Not legal advice; a local lawyer reviews before any live outbound.

ID	Requirement	Phase	Acceptance
CMP-01	Per-contact lawful basis, data source, consent timestamp	D	Contact without a basis is held for gated channels
CMP-02	Country channel rules as data: KSA and Egypt consent or documented relationship for marketing; Egypt marketing-licence field; UAE consent-first; EU blocked by default	D	Blocked send shows the rule
CMP-03	Opt-out in every message; global suppression that survives re-import	D	Suppressed contact stays suppressed
CMP-04	Data-residency flag per workspace	D+	Storage region recorded
CMP-05	Data-subject export and delete across records, assets, caches	N	Deletion test removes linked data and blocks re-import

7.11 Billing and metering

ID	Requirement	Phase	Acceptance
BILL-01 *(new)*	Plan limits as data (§4.2): seats, records, AI credits, enrichment credits, sends, exports, storage	D	Changing a limit needs no code change
BILL-02 *(new)*	Usage meter per action with a published cost table in the app	D	Each spend writes a ledger entry with action, units, plan and cost
BILL-03 *(new)*	Spend controls: hard stop on Free; workspace cap on by default; per-member caps; alerts at 75/90/100%	D	No run exceeds a cap in tests
BILL-04 *(new)*	Plans page and mock checkout for the demo	D	Clearly labelled as not live
BILL-05 *(new)*	Real billing through a merchant of record, then local gateways (Paymob, Moyasar/HyperPay); prepaid packs; annual plans	L	Tax and seller eligibility confirmed in writing first
BILL-06 *(new)*	Agency client rebilling and white-label reports	L	Client invoice matches usage ledger
BILL-07 *(new)*	Free-workspace archive after 60 days idle, one-click restore	N	Restore returns all data

7.12 Founder content (second track)

ID	Requirement	Phase	Acceptance
CON-01	Brand voice profile from the founder's own posts and website	D+	Cites its examples
CON-02	Post drafts in the founder's voice and language, human-approved	N	Same Arabic checks as ARB-03
CON-03	Scheduling and posting via `w_member_social`	N	Founder OAuth; limits respected
CON-04	Content calendar and performance import	L	Metrics carry source and date

7.13 Later modules (full catalogue)

Every v1 module stays in the product roadmap. Detailed rows and acceptance criteria are in v1 §5; their v3 phases are in Appendix A.

Module	v1 IDs	v3 phase
CRM core (workspaces, companies, people, tasks, dedupe, exclusions, roles)	CRM-01–10, UI-01, UI-03	D for what plays need; N for pipelines and custom objects
Sales (qualification, deals, meeting prep, proposals, quotes, forecasts, renewals)	SALE-01–12	N (SALE-01, 03, 05, 06, 07, 09), L (rest)
Marketing (briefs, segments, nurture, social, landing pages, events, paid, attribution)	MKT-01–15	N (MKT-04, 06, 07), L (rest)
Integrations (external CRMs, Make/webhooks, forms, GA4, Clarity)	INT-06–16	N (INT-06, 11–13), L (rest)

Module	v1 IDs	v3 phase
Operations (rules, review queues, jobs, backups, client portal, monitoring)	OPS-01–09	D (OPS-03, 04, 06, 09), N (OPS-01, 02, 05, 07), L (OPS-08)
AI evaluation (ranking, retrieval, assistant safety, adaptation, commercial experiments)	EVAL-01–05	D (EVAL-03), N (EVAL-02, ARB-05), L (EVAL-01, 04, 05)

7.14 From the real-workflow teardown *(v3.1)*

Source: research 11. Each row replaces a manual step seen in a real engagement.

ID	Requirement	Phase	Acceptance
GAP-01	Single source of truth with versions: imports diff against the previous version instead of re-exported copies	D+	Re-importing an edited file shows added, changed and unchanged rows; no duplicate records
GAP-02	CRM duplicate check before import (domain, email, name, alias) with a merge preview	D	Near-duplicates in the demo data are flagged before import
GAP-03	Triage buckets beyond score tiers: client, partner, drop, existing client, with rationale and owner; Keep/Drop review queue	D	Existing clients never enter cold sequences; partner accounts route to P17
GAP-04	Angle engine: segment × persona × HR-maturity (P1a/P1b/P2) → offer angle; register (founder or leader voice) from the recipient's title	D	Draft shows which angle and register it used; changing the title changes the register
GAP-05	Native sender push replaces Make.com glue: approved copy and lead list go to the sender platform in one step	D (sandbox)	No external automation tool needed for the demo flow
GAP-06	Launch gate per campaign: a checklist of blockers (missing research, missing basis, unverified emails, sender pre-flight) blocks queueing	D	Queueing is refused until every item is cleared or overridden with a note
GAP-07	Funnel scorecard from sequence events: connect → accepted → reply → meeting → agreed, written back to the CRM	D+	Scorecard numbers match the event log
GAP-08	Call outcome logging by voice: the SDR dictates, Whisper transcribes, the assistant proposes a structured outcome for confirmation	D	Arabic and English notes produce a confirmed outcome on the record
GAP-09	Reference library: social-proof references with a consent flag, rotated automatically per contact	D+	A reference without consent is never inserted
GAP-10	Message confidence score from evidence count, freshness, register fit and missing fields	D+	Low-confidence drafts are labelled before review
GAP-11	SDR onboarding kit generated from ICP, offers, plays and rubric (extends PB-06)	D+	A new member receives a PDF kit on joining a workspace

ID	Requirement	Phase	Acceptance
GAP-12	Discovery-pilot capture (play P15): interview answers, pain scores, maturity score, stored per account	D+	Scores aggregate per segment
GAP-13	Owner assignment per contact with round-robin and per-member caps (extends GOV-05)	D	Seeded Team workspace shows three SDR owners
GAP-14	Event-claim rule: event references stay at company level ("your company was at…"), never personal	D	ARB-03 style check flags personal event claims

7.15 AI harness and voice *(v3.1)*

ID	Requirement	Phase	Acceptance
AI-01	Own agent loop with streaming, typed tools from `packages/contracts`, step cap 8; spending or sending tools pause for approval	D	Live run streams tokens; a spending tool produces a quote, never a run
AI-02	Prompts as versioned files; prompt version recorded on every run and chat message	D	Changing a prompt changes the recorded version
AI-03	Provider adapter: Azure OpenAI now, OpenAI-compatible MaaS later, fixture for tests	D	Switching provider is configuration only
AI-04	Voice dictation in the assistant and call outcome forms (≤ 60 s), Whisper transcription with an en/ar hint, credits quoted	D	Arabic and English clips transcribe; audio is not stored by default
AI-05	Agent evaluation set: tool choice, refusal to spend without approval, Arabic quality rubric; transcription WER on a small clip set	D+	Results recorded with date and model
AI-06	Twenty-side chat can call our product MCP so plays run from inside the CRM	D+	A play started in Twenty appears in the review queue

8. Architecture

8.1 Components

Layer	Component	Licence	Decision
CRM and data model	Twenty	AGPLv3 with the Twenty Application Exception; SDK and app toolkit MIT; files marked `@license Enterprise` are commercial	Build as a Twenty app or separate service on its API/SDK. **Do not modify or fork Twenty**: changes to Twenty itself must be published under AGPL
Asset Studio	OpenDesign 0.24.x	Apache-2.0	Integrate unmodified behind our wrapper; add our own skills

Layer	Component	Licence	Decision
Product services	Play engine, waterfall, sequence state, compliance, metering, product MCP server	Ours	Build
Automation connectors	Activepieces pieces	MIT outside EE folders	Reuse pieces later (N); confirm the EE boundary first. **Not n8n**: its licence allows internal business use only
Public-web research	Crawl4AI	Apache-2.0	Integrate
Email verification fallback	Reacher	MIT	Integrate (D+)
LLM tracing	Langfuse	MIT core	Integrate (D+)
Link and page analytics	Umami	MIT	Integrate (N)
Inbox and WhatsApp	Chatwoot	MIT core	Integrate (N)
Email sending	Instantly or Smartlead (customer's mailboxes)	Commercial	Integrate via API; we do not run shared sending IPs

Do not use: n8n as an embedded engine; Windmill, Firecrawl, Plausible, Listmonk, DocuSeal or Documenso as modified hosted forks (AGPL); any scraper or auto-updating (`@latest`) code with access to customer credentials; shared-IP sending or warm-up networks; open-source "Clay clones" as a base. Reasons: research 09.

8.2 Azure placeholder (until December)

Owner update, 14 October: keep the application on localhost while testing. Prepare one cost-efficient Linux server with Terraform, a separate project resource group and a Reader-only GitHub connection. No staging/production split, availability zones, application deployment or model provisioning yet. The managed services below remain an expansion option; the current target uses Docker PostgreSQL/Redis and persistent local file/vault adapters on the server disk. See Azure preparation.

Function	Azure service
LLM for assistant, research and drafts	Azure AI Foundry model endpoint behind our OpenAI-compatible gateway
App, workers, Twenty, OpenDesign	Azure Container Apps or AKS, same images as development
Database and cache	Azure Database for PostgreSQL, Azure Cache for Redis

Function	Azure service
Files and assets	Azure Blob Storage (Twenty S3-compatible storage settings through an adapter)
Secrets	Azure Key Vault

8.3 Huawei Cloud (after passing December)

We receive Huawei Cloud resources only if we pass the December phase. Then every row below replaces its Azure placeholder.

Function	Huawei service	Portability rule
LLM	ModelArts Studio (MaaS): DeepSeek, Qwen or GLM through the same OpenAI-compatible gateway	Provider switch is one configuration change; Arabic evaluation reruns on the new model
Compute	CCE or ECS	Same container images
Database and cache	RDS for PostgreSQL, DCS (Redis)	Backup and restore test before cut-over
Files	OBS	S3-compatible adapter
Secrets	DEW / KMS	Same vault interface
Evaluation and any fine-tuning	ModelArts notebooks and jobs	Results record model, region and date

The December architecture plan shows the Azure placeholder, the Huawei target and the cut-over steps from v1 §8.2. Every provider-specific service sits behind an adapter (model gateway, S3-compatible storage, Postgres, Redis, vault) so the switch is configuration, not code.

9. Moat and pushback

The product owner asked for pushback. Points where the evidence disagrees, and what we do about each:

1. **Do we need LinkedIn? Yes as a channel, no as a scraper.** LinkedIn is where B2B buyers in Egypt and the Gulf are reachable, and Lara's sequence uses it twice (invite and follow-up). The product needs three LinkedIn things: (a) manual steps with copy-ready text, (b) posting from the founder's own profile through the official API (plays P6, P7), and (c) people and company data, which licensed providers (Crustdata, Apollo, etc.) and customer exports already give without touching LinkedIn. None of these needs crawling. Agent-Reach routes LinkedIn to `stickerdaniel/linkedin-mcp-server`, which drives the user's logged-in browser and can send messages; its own FAQ says accounts "can be restricted or banned". Run on our servers with customers' cookies, it repeats the hiQ, Proxycurl and ProAPIs fact pattern and adds PDPL duties. So: nothing LinkedIn-automated before December; when we need deeper LinkedIn data, compare the lower-risk options in §7.7 (licensed providers first, then a local one-profile-per-click connector).
2. **"Gather in one place" is not a moat by itself.** GoHighLevel owns "all-in-one", and every vendor ships MCP connectors. Salesforce's own survey says 84% of teams want to consolidate, so incumbents chase the same demand. Our moat must be what accumulates: the play library tuned for MENA, Arabic evaluation data, compliance rules as data, asset templates that work in Arabic, and reply outcomes per play.

3. **Design assets are a bet.** No independent study shows that one-pagers, audits or carousels raise B2B replies; the only side-by-side test found was small, vendor-run and mixed. Keep the Asset Studio as a major pillar (D6), but ship every asset play with an asset-vs-no-asset holdout and lead the pitch with the stronger pains (first campaign without learning a tool, cost control, safe setup).
4. **Free tiers cost money and convert poorly.** The best public figure is a median 8% free-to-paid (self-reported, vendor co-produced); plan for 2–4%. The free tier must hard-stop at quota and cost under about $2–3 per active user per month worst case.
5. **Payments are hard from Egypt and KSA.** Stripe does not list either country. Billing is mocked in the demo; seller eligibility must be confirmed before launch.
6. **Licences can bite.** Twenty is AGPL: build on its API/SDK, never fork it. n8n cannot be embedded in a multi-tenant SaaS.
7. **English default lowers the Arabic advantage in the competition.** Show Arabic as a toggle in the demo (one Egyptian draft and one Arabic asset), and keep the Arabic evaluation for phase N and the MENA pitch.
8. **Cold outbound is legally constrained.** Egypt enforces its PDPL from 31 October 2026, with a licence and consent for electronic direct marketing; Saudi Article 28 already requires consent for marketing to people with no prior interaction. Plays lead with warm paths, lead magnets and consent; cold email is one play type behind country gates.
9. **Full catalogue vs focus.** Keeping every feature in the catalogue is fine; building them in parallel is not. Phase D stays the outbound loop with four plays and two templates; pull from D+ only when D passes its acceptance tests.

10. Success measures

Measure	Target for the demo
Time from empty workspace to first approved drafts with play P1	≤ 30 minutes, recorded rehearsal
Non-technical tester connects mailbox and one provider	≤ 10 minutes without documentation
Asset from template to approved PDF	≤ 5 minutes
Unapproved sends, duplicate contacts, runs over cap	0 (blocking)
Arabic draft acceptance by native reviewers	Reported per dialect in phase N
Reply per delivered (real, consented pilots only)	Reported with sample size; no lift claim from fixtures

11. Plan

The December deadline is a date in December; the plan aims to finish by 30 November.

Window	Work	Exit gate
13–19 Oct	Assistant spike (AST-01); Twenty + OpenDesign running on Azure; data model for plays, signals, assets and the usage ledger; 5–10 founder/agency interviews start	Spike decision written; both engines run on Azure
20 Oct – 9 Nov	ONB-01/02, plays P1–P2, scoring, waterfall (fixtures + one sandbox), research, English/Egyptian drafts, approvals, compliance gates, metering and caps; idea PDF first draft	Play P1 runs end to end on fixtures

Window	Work	Exit gate
10–30 Nov	Plays P3–P4, Asset Studio (STU-01–03, 07–08), sequence via sandbox sender, LinkedIn/call tasks, audit, plans page	Demo website on Azure shows one play and two assets with gates
1 Dec – deadline	Idea PDF final, architecture plan, demo polish and recording; D+ items if time	PDF sent
If we pass	Huawei migration (§8.3), phase N plays and templates, agency mode, Arabic evaluation, WhatsApp opt-in, public repository (D17)	Regional stage

12. Open questions

Gaps found in v3.1 (need decisions): (a) data residency for voice audio and transcripts, especially for KSA residency workspaces; (b) Arabic transcription accuracy per dialect has no evaluation yet (only a small clip set); (c) prompt injection from scraped research now reaches a tool-calling loop, so injection tests (RES-03) must run against the live agent; (d) AI cost per play is unknown until live runs are logged; (e) agency parent/client relationships are stored in workspace context for the demo, not as a first-class model; (f) demo workspaces share the two local Twenty workspaces, so CRM isolation is only proven for Revenue Demo A/B.

1. Interview list: which 5–10 founders and agencies (Cairo, Riyadh/Dubai) can we talk to before the PDF is written?
2. Which sender (Instantly or Smartlead) and which signal provider (PredictLeads or Crustdata) gives us a sandbox?
3. Who reviews Gulf Arabic in phase N?
4. Is a legal opinion on Egypt PDPL B2B direct marketing available (university, ITIDA or a mentor) before any live pilot?

13. Sources

Research reports: 01 GTM process and AI, 02 competitors, 03 connectors, 04 MENA regulation and Arabic AI, 05 Huawei competition and cloud, 06 Huawei 2026–27 rules, 07 pricing and freemium, 08 market data, 09 building blocks and Agent-Reach, 10 user pain and out-of-the-box. Lead-reviewer spot checks: research README. Meeting: `Lara Meeting transcript.txt`. v3.1: 11 real workflow teardown.

Appendix A. v1 feature catalogue with v3 phases

v1 ID	Feature	v3 phase / v3 ID
CRM-01	Agency/client workspaces	D (single workspace), N (GOV-08 agency)
CRM-02	Companies and contacts with relationships	D
CRM-03	Lead ownership and routing	D (GOV-05), N (routing rules)
CRM-04	Opportunity pipelines	N
CRM-05	Tasks, reminders, notes, timelines	D

v1 ID	Feature	v3 phase / v3 ID
CRM-06	Tags, segments, saved views, bulk selection	D (segments via PB-01), N (rest)
CRM-07	Validation, deduplication, reviewed merge	D
CRM-08	Relationship history, existing clients, exclusions	D
CRM-09	Custom objects and fields	D (product objects), N (user-defined)
CRM-10	Roles, record access, client sharing	D (GOV-06), N (client sharing)
UI-01	Twenty-derived design system	D
UI-02	English-only experience	Replaced by D1 (English default, Arabic content, Arabic UI later)
UI-03	Loading, empty, error, partial states	D
AI-01–02	Persistent assistant and live context	D (subject to AST-01)
AI-03	Cited workspace knowledge	D
AI-04	In-app navigation	D
AI-05	Typed actions	D (GOV-03)
AI-06	Multi-step plans and batch execution	D (plays)
AI-07	Proactive next-action suggestions	N
AI-08	Client knowledge and campaign memory	D (play versions), N (memory)
AI-09	Design mode with OpenDesign	D (AST-04, STU)
AI-10	Durable jobs, receipts, cancellation	D
AI-11	Tool policies, approval scopes, spend limits	D (BILL-03, GOV)
AI-12	Stale context and untrusted instructions	D (RES-03)
AI-13	Model and prompt versions, usage observability	D+ (GOV-09)
AI-14	Plain-English rule editing	L
AI-15	Analytics explanations	L
GTM-01	Import and reconciliation	D (SIG-01)
GTM-02	Offer catalogue, ICP, proof	D (PB-01, PB-03)
GTM-03	Document/playbook upload and retrieval	D+
GTM-04	Research with citations	D (RES-01)
GTM-05	Account–role–offer ranking	D (SCR-01)
GTM-06	ABM plans and account briefs	D (DES-02), N
GTM-07	Enrichment with provenance and budget	D (ENR)
GTM-08	Personalized outreach and voice	D (MSG-01), D+ (MSG-02)
GTM-09	Sequence branches and schedules	D (SEQ-01)

v1 ID	Feature	v3 phase / v3 ID
GTM-10	Campaign versioning and approval invalidation	D (PB-04)
GTM-11	Campaign QA, collisions, exclusions	D (GOV-05, CMP-03)
GTM-12	Sender readiness and holds	D (SEQ-03)
GTM-13	Export and readiness views	D
GTM-14	Partnerships	L
GTM-15	Event attendee follow-up	N
GTM-16	Research refresh and evidence expiry	N
SALE-01	Qualification and conversion	N
SALE-02	Deal stages and ownership	N
SALE-03	Next steps and overdue follow-up	N
SALE-04	Interaction history and handoff	N
SALE-05	Communications view and reply triage	D+ (SEQ-06), N
SALE-06	Meeting preparation and summaries	N (play P10)
SALE-07	Proposals, decks, case studies	N (STU-04/05)
SALE-08	Basic quotes	L
SALE-09	Won/lost and handoff	N
SALE-10	Pipeline health and forecasts	L
SALE-11	Renewals and expansion	L
SALE-12	Team playbooks	D (plays)
MKT-01	Campaign briefs	L
MKT-02	Segments and lists	N
MKT-03	Multi-channel calendar	L
MKT-04	Content briefs and copy	N (CON-02)
MKT-05	Nurture and lifecycle email	L
MKT-06	Social planning and publishing	N (play P7, CON-03)
MKT-07	Landing pages and lead capture	N (STU-04)
MKT-08	Events and webinars	L
MKT-09	Partner and referral campaigns	L
MKT-10	Paid campaign planning	L
MKT-11	Paid/social publishing adapters	L
MKT-12	UTM attribution and funnel reporting	N (Umami), L

v1 ID	Feature	v3 phase / v3 ID
MKT-13	Experiments and variants	D+ (SEQ-08)
MKT-14	Marketing operations dashboard	L
MKT-15	Measurement plans	L
DES-01	Brand profiles and design-system packages	D (STU-02)
DES-02	Embedded Design mode	D (AST-04)
DES-03	Account research brief	D (DES-02 v2)
DES-04	ABM one-pager	D (STU-03)
DES-05	Campaign summary and client report	N (play P13, STU-05)
DES-06	Networking card with QR	L
DES-07	Proposals and sales collateral	N (STU-05)
DES-08	Social, ad, email, landing creatives	D+ / N (STU-04/05)
DES-09	Conversational revisions and lineage	D+ (STU-06)
DES-10	Exports and asset library	D (STU-07/08)
DES-11	Brand, accessibility and chart checks	N
INT-01–02	Smartlead preparation and launch	D (SEQ-02, sandbox), N (live with consent)
INT-03	Dripify automation	Dropped (D7)
INT-04	Call tasks	D (SEQ-05)
INT-05	WhatsApp	N (SEQ-07)
INT-06	Email and calendar sync	N
INT-07	Enrichment providers	D (ENR)
INT-08	External CRM connectors	L
INT-09	Make/webhooks	L (Activepieces pieces)
INT-10	Forms, events, social/ad connectors	L
INT-11	Provider-event inbox and reconciliation	D (SEQ-02)
INT-12	Credential health and rate limits	D (GOV-01, GOV-04)
INT-13	Connector catalogue	D (GOV-01)
INT-14	GA4	L
INT-15	Microsoft Clarity	L
INT-16	Cross-source identifiers	L
OPS-01	Versioned rules	N
OPS-02	Review queues and notifications	D

v1 ID	Feature	v3 phase / v3 ID
OPS-03	Durable jobs and idempotency	D
OPS-04	Audit log	D (GOV-07)
OPS-05	Settings and administration	D
OPS-06	Usage, budgets, cost reporting	D (BILL-02/03)
OPS-07	Backups, restore, deletion	N (CMP-05)
OPS-08	Client reporting and portal	N (GOV-08), L (portal)
OPS-09	Monitoring and incident recovery	D
EVAL-01	Ranking experiment	L
EVAL-02	Retrieval and grounded generation	N
EVAL-03	Assistant safety tests	D
EVAL-04	Model adaptation	L
EVAL-05	Commercial experiments	L (needs real outcomes)

Changelog

v3.1 — 15 October 2026. Added decisions D18–D21 (Raycast-derived design system, own AI harness on Vercel AI SDK Core, Azure models and Whisper approved, real data read for structure only). Added plays P15–P17, requirement sections §7.14 (14 features from the real-workflow teardown) and §7.15 (AI harness and voice), and the v3.1 gaps in §12. ARB-07 (Arabic UI) is partly delivered early: interface strings and RTL layout exist; Twenty's own RTL is unchanged. Research 11 added.

v3 — 13 October 2026. Owner decisions on the v2 proposals, business model, plays, Asset Studio, architecture and plan.

Implementation checkpoint — 15 October 2026 continuation

The owner's later unattended-run instruction authorizes project-only Azure deployment and inference. Repository pushes and live outreach remain outside that authorization.

Requirements	Implemented behavior	Remaining acceptance work
GAP-01	Re-import reports added/updated/unchanged outcomes and invalidates stale reviews	File version history UI
GAP-02	Near-duplicate import holds, workspace review, atomic exclusion and sequence cancellation	CRM field-merge preview
GAP-03/04/14	Buckets/rationale/owner/segment controls, angle/register drafting, personal event-claim checks at approval and queue boundaries	Partner play P17 is still later scope
GAP-06	Visible launch checklist; queue-time holds remain authoritative	A note never bypasses suppression or consent gates

Requirements	Implemented behavior	Remaining acceptance work
GAP-07	Event scorecard and asset/holdout counts; confirmed call outcomes sync to CRM	LinkedIn connect/accepted telemetry and general scorecard CRM writeback
GAP-09/10	Consent-managed reference library, reference rotation and visible evidence-based draft confidence	Native reviewer calibration of confidence
GAP-11	Branded SDR kit PDF from current workspace context	Automatic delivery on member joining
GAP-12	Account discovery capture and segment pain/maturity aggregates	Bilingual interview guide/questionnaire generation
GAP-13	Manual member assignment; opt-in balanced import distribution with per-member contact caps	Contact caps currently apply to automatic import assignment

These are implementation states, not claims of external-provider activation or full PRD completion. See validation for executed checks and limitations.